

Adaptive Memory Selection for Long-Term AI Agents Using Machine Learning

Raghavendra Beri (002591668)

Harsha Prakash (002078092)

MS Data Science, Northeastern University

CS6140: Machine Learning — Final Project Report

Abstract

Long-term LLM agents accumulate more conversational history than fits in a context window, forcing a choice about which past “memories” to recall. The prevailing solution — embedding-similarity (top- K) retrieval — surfaces memories that merely *sound like* the query, not those that will actually be *used* again. We reframe recall as supervised binary classification: from sixteen interpretable features computed the moment a new session begins, predict whether each past utterance will be referenced later (`retain`) or not. Using 17,631 memories from 450 Multi-Session Chat (MSC) conversations, we compare Logistic Regression, k -NN, Decision Trees, Random Forest, and AdaBoost against a tuned similarity-threshold baseline. All five classifiers beat the baseline substantially (ROC-AUC ≈ 0.76 vs. 0.63; F1 ≈ 0.43 vs. 0.34), and feature analysis shows the gain comes from *specificity* and *persona-salience* signals that similarity ignores. Integrated with a local Llama 3.1 8B model, classifier-selected memories match store-everything answer quality using roughly one-quarter of the context tokens.

1 Introduction

Conversational AI agents are increasingly expected to persist across days and hundreds of turns, remembering a user’s stated facts, preferences, and history. Because a transformer’s context window is bounded, an agent cannot re-read its entire past at every turn; it must *select* a small subset of memories to condition on. The dominant industrial practice stores each past utterance as a vector and, at query time, retrieves the top- K most cosine-similar ones. This is cheap but blunt: similarity to the current phrasing is a poor proxy for future usefulness. A throwaway greeting can be highly similar to a later greeting, while a durable fact (“I work as a stunt double”) may be phrased nothing like the question that eventually depends on it.

We ask whether supervised machine learning can predict memory importance more effectively than similarity retrieval. We treat each memory as a training example and define a binary **target**, `label_retain`: it is 1 when the memory is explicitly referenced or paraphrased in a later turn of the conversation, and 0 otherwise. This target operationalizes “will this memory be useful again?” in a way that is measurable directly from dialogue and, crucially, is *not* defined by embedding similarity — so a similarity-based retriever cannot be optimal by construction.

Three properties make the task hard. First, **no gold relevance labels exist**: MSC was not annotated for memory usefulness, so labels must be induced from the data with a reference heuristic and then validated. Second, the classes are **highly imbalanced** — only 15.6% of memories are ever referenced again — so naive accuracy is misleading and the minority (`retain`) class is easy to miss. Third, the inputs are **short, noisy, first-person utterances** whose importance depends on subtle cues (named entities, specificity, self-disclosure) rather than surface form. Sections 5–8 describe how we address each.

2 Proposed Method

We engineer sixteen features that are all computable at *recall time* (the start of a new session) — temporal (recency, recurrence), structural (speaker, position), lexical (length, TF-IDF specificity, stopword ratio, entities), and semantic (SBERT cosine to the opening query, and salience against the agent’s carried-forward persona summary). We standardize them, correct class imbalance with SMOTE, and train five classifiers under grouped, stratified cross-validation, selecting the best by F1 on held-out conversations. We believe this beats top- K similarity retrieval — the approach we most directly compete with — because similarity captures only one of our sixteen signals, and our own analysis (Sec. 9) shows that signal is far from the most predictive: information-density and persona-salience features carry more of the signal about what gets reused, and a learned combination of features dominates any single one.

3 Related Work

Long-term dialogue and memory. Xu, Szlam, and Weston [1] introduced the Multi-Session Chat dataset and the “Beyond Goldfish Memory” setting we build on; they showed that summarizing prior sessions into a persona memory and retrieving from it improves engagingness, using a retrieval-augmented generator with a learned retriever. Our work differs by treating memory *selection* itself as an explicit, interpretable classification problem rather than folding it into an end-to-end neural retriever. Packer et al.’s MemGPT [3] manages a tiered memory hierarchy with the LLM itself deciding what to page in and out via function calls; this is powerful but opaque and token-expensive, whereas our classifier is a lightweight, auditable gate. Zhong et al.’s MemoryBank [4] adds an Ebbinghaus-style forgetting curve that decays memories by time and rehearsal — essentially a hand-designed function of our `recency` and `access_frequency` features; we instead *learn* how such factors combine, and find recency is a comparatively weak predictor.

Sentence representations. Our semantic features use Sentence-BERT embeddings [2], which fine-tunes BERT with a siamese objective so that cosine distance is semantically meaningful; we use the compact `all-MiniLM-L6-v2` variant. Reimers and Gurevych’s key finding — that mean-pooled siamese embeddings drastically outperform raw BERT for similarity — is what makes our `sim_to_current_query` and `persona_salience` features usable, though we corroborate that similarity alone is insufficient for importance ranking.

Learning algorithms. We rely on established classifiers. Breiman’s Random Forests [5] average de-correlated bagged trees to reduce variance and yield feature importances, which suits our mixed-scale, interaction-heavy features. Freund and Schapire’s AdaBoost [6] sequentially reweights hard examples; it is a natural complement on imbalanced data. Chawla et al.’s SMOTE [7] synthesizes minority-class examples along feature-space segments between neighbors, and is our primary imbalance remedy; we validate its recall benefit empirically. Together these citations motivate both our feature design ([1, 2, 4]) and our modeling and imbalance choices ([5, 6, 7]); relative to all of them, our contribution is an explicit, cheap, and interpretable memory-selection classifier that we show outperforms the similarity retrieval those systems typically rely on.

4 Related Implementations (Kaggle)

The MSC corpus is a research dataset distributed through ParlAI and is *not* hosted on Kaggle, so no Kaggle kernel targets it directly. We therefore draw on Kaggle work for the two analogous sub-problems our task combines: short-text binary classification and imbalanced learning.

The “Natural Language Processing with Disaster Tweets” [8] getting-started competition is the closest text analogue: contestants classify short, noisy user utterances as relevant/not, and the dominant public kernels pair TF-IDF or embedding features with Logistic Regression and linear SVMs — the same specificity-plus-linear-model recipe that our mutual-information analysis independently favors. Our task is harder in that positives are

defined by cross-turn reference rather than a fixed topic, and we add temporal and persona features those kernels lack.

The “Credit Card Fraud Detection” dataset [9] is the canonical Kaggle playground for extreme class imbalance; its most-upvoted notebooks demonstrate that accuracy is meaningless at low prevalence and that SMOTE combined with tree ensembles and PR-AUC evaluation is the standard remedy. We adopt exactly this evaluation discipline (PR curves, PR-AUC, SMOTE ablation) but on text-derived features rather than anonymized PCA components. Finally, the “Toxic Comment Classification” [10] challenge popularized treating nuanced text-relevance as supervised classification with heavy class skew; like those kernels we lean on TF-IDF specificity, but we show that adding a learned persona-salience feature — unavailable in a single-message setting — yields signal beyond bag-of-words. Across all three, our method is better suited to *conversational* memory because it encodes the temporal and cross-session structure that single-message Kaggle tasks do not have.

5 Dataset and Data Analysis

Source and construction. We use Multi-Session Chat (MSC v0.1) [1], downloaded via ParlAI. Each episode in session N contains the full transcripts of sessions $1:N-1$ (the `previous_dialogs` field) plus the current session and a carried-forward persona summary. We treat every utterance in the previous sessions as one *memory*, the opening turn of the current session as the *current query*, and the remaining current-session turns as the *future* against which reference is measured. Parsing the validation splits of sessions 3–5 (150 episodes each) yields **17,631 memories across 450 episodes**, averaging 39 memories per conversation.

Label induction and validation. Because MSC has no memory-relevance labels, we induce them (module `src/labeling.py`, notebook 01). A memory is `retain=1` if any later turn (excluding the opening query, which feeds a feature) satisfies one of three triggers: (i) a shared proper noun or number (*named* reference); (ii) at least two shared *specific* content words, where “specific” means document frequency below 3%; or (iii) paraphrase, i.e. SBERT cosine > 0.60 to the memory. This yields a 15.6% positive rate (1,778 *content*, 558 *named*, 421 *paraphrase* triggers). We validate these automatic labels against MSC’s *human* per-utterance persona-worthiness annotation (recovered for 100% of memories from the `msc_personasummary` split). Persona-worthy memories are referenced later significantly more often than non-worthy ones (18.5% vs. 10.6%; odds ratio 1.92; $\chi^2 = 191.9$, $p < 10^{-43}$). At $n = 17,631$ statistical significance is nearly automatic, so it is the odds ratio (1.92), not the vanishing p -value, that carries the practical meaning: a persona-worthy memory is about twice as likely to be re-referenced. This confirms the label captures genuine importance structure — while the modest agreement (Cohen’s $\kappa = 0.06$) shows that “referenced-again” and “human-persona-worthy” are related but distinct constructs, an insight we return to in Section 9.

Features and formulas. Each memory is embedded with SBERT `all-MiniLM-L6-v2` (384-d) [2]. The semantic features use cosine similarity,

$$\cos(\mathbf{u}, \mathbf{v}) = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}, \quad (1)$$

computed as `sim_to_current_query` (memory vs. opening query) and `persona_salience` (max cosine to the carried-forward persona sentences). Lexical specificity uses TF-IDF, $\text{tfidf}(t, d) = \text{tf}(t, d) \cdot \log \frac{N}{1+\text{df}(t)}$, summarized per memory by its max and mean nonzero weight. `access_frequency` counts earlier same-episode memories with cosine > 0.6 (a recurring topic). The remaining features are recency in sessions, turns since the memory was spoken, position in its session, speaker, sentence and character length, stopword ratio, and binary flags for numbers, proper nouns, questions, and first-person mentions. All are available at recall time; none use future turns, preventing leakage into the label.

Class balance and discriminative features. Figure 1 confirms the 15.6%/84.4% split. Figure 1 shows that retained memories are longer, more information-dense (higher TF-IDF), and markedly more persona-salient than discarded ones, whereas `sim_to_current_query` separates the classes only weakly.

Redundancy and information content. The correlation matrix (Fig. 5, Appendix) shows expected clusters




Figure 1: (a, left) The target is imbalanced. (b, right) Retained memories differ most in specificity and persona-salience, less in raw query similarity.




Figure 2: (a, left) Specificity and persona features are most informative; `sim_to_current_query` ranks only seventh. (b, right) The feature space needs ~ 10 PCA components for 90% variance.

(sentence and character length; the two TF-IDF summaries) but no feature is collinear with the target, motivating a multivariate model. Ranking features by mutual information with the label, $I(X; Y) = \sum p(x, y) \log \frac{p(x, y)}{p(x)p(y)}$, the top signals are `tfidf_mean` (0.053), `tfidf_max` (0.053), `persona_salience` (0.045), and `char_length` (0.035); `sim_to_current_query` ranks only seventh (0.020) (Fig. 2). This is the quantitative heart of our argument: the feature that similarity retrieval relies on is not the most informative one.

Dimensionality. Standardizing and running PCA (Fig. 2), about ten components are needed to explain 90% of variance, and the first two components do not linearly separate the classes — consistent with the nonlinear models (Random Forest, AdaBoost) performing competitively and confirming that no small linear subspace trivially solves the task.

6 Learning Algorithms

All models are trained inside an imbalanced-learn pipeline `StandardScaler` $\rightarrow$ `SMOTE` $\rightarrow$ classifier, with `SMOTE` applied only to training folds.

Logistic Regression models $P(y=1 \mid \mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x} + b)$ with $\sigma(z) = 1/(1 + e^{-z})$, fit by minimizing

regularized log-loss $\sum_i \log(1 + e^{-y_i(\mathbf{w}^\top \mathbf{x}_i + b)}) + \frac{1}{2C} \|\mathbf{w}\|^2$. It is linear and highly interpretable — its standardized coefficients (Fig. 4) directly rank feature influence — and serves as our transparent reference model.

k-Nearest Neighbors is non-parametric: a memory is scored by the (distance-weighted) label fraction of its k closest training points in standardized feature space. It captures local structure with no training-time fitting, but degrades in higher dimensions and is sensitive to feature scaling, which is why standardization matters here.

Decision Trees recursively split features to reduce impurity; using Gini impurity $G = 1 - \sum_c p_c^2$, each split maximizes the weighted impurity decrease. A single tree is readable and models feature interactions and nonlinearity, but overfits, so we bound depth and leaf size via cross-validation.

Random Forest [5] averages many trees, each grown on a bootstrap sample using a random feature subset per split; this decorrelation reduces variance, $\text{Var} \approx \rho \sigma^2 + \frac{1-\rho}{B} \sigma^2$, and gives robust out-of-the-box performance plus impurity-based feature importances (Fig. 4). It suits our mixed-scale, interaction-rich features.

AdaBoost [6] fits a weighted sum of shallow trees $H(\mathbf{x}) = \text{sign}(\sum_t \alpha_t h_t(\mathbf{x}))$, at each round up-weighting misclassified examples and setting $\alpha_t = \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$. By focusing on hard cases it can sharpen the minority class, complementing the bagged forest.

Hyperparameters (C ; k and weighting; tree depth and leaf size; number and learning rate of boosting rounds) are tuned by grid search under the cross-validation described next.

7 Experimental Setup

The full pipeline is reproducible from `src/` (see the Appendix) and summarized in notebook 03. We use all 17,631 parsed memories. To prevent leakage between conversations, we split at the *episode* level with `GroupShuffleSplit` (80%/20%), so every memory from a conversation lands entirely in train or test: 14,146 memories (360 episodes) for training and 3,485 (90 episodes) for testing.

Model selection uses 5-fold `StratifiedGroupKFold` on the training set: folds preserve the 15.6% positive rate *and* keep episodes intact, so cross-validation never scores a model on a conversation it partly trained on. Grid search optimizes F1 (the operative metric under imbalance). Critically, the *single deployed model is chosen by mean cross-validated F1* — the test set is never consulted for model choice. Only after that choice is fixed do we evaluate each model once on the untouched test episodes, reporting accuracy, precision, recall, F1, ROC-AUC, and PR-AUC (average precision). This keeps the headline number an honest generalization estimate rather than a maximum taken over models on the test set.

Two controls make the comparison meaningful. The **similarity baseline** stands in for top- K retrieval: we tune a single threshold on `sim_to_current_query` to maximize training F1, then apply it to the test set — isolating exactly the signal an embedding retriever uses. Separately, an **imbalance ablation** refits the best model type under three regimes — no resampling, class-weighting, and SMOTE — to quantify what imbalance handling buys. Randomness is fixed (`seed=42`) throughout.

8 Results

Table 1 reports held-out test performance. Every supervised model outperforms the similarity baseline by a wide margin: the deployed Logistic Regression lifts ROC-AUC from 0.633 to 0.767 and F1 from 0.339 to 0.429 — a relative F1 improvement of 27% (and every other model clears the baseline too). The ROC curves (Fig. 3) make the gap visual: all five learned models dominate the dashed baseline across the operating range, while the baseline sits only modestly above chance.

Which model. We choose the deployed model strictly by mean cross-validated F1 on the training folds; the test set plays no part in the choice. By that criterion **Logistic Regression wins** (CV F1 0.426, CV ROC-AUC 0.777), narrowly ahead of AdaBoost (0.421) and Random Forest (0.406), and its held-out scores (F1 0.429,

Table 1: Held-out test metrics (best per column in **bold**). The similarity baseline is the incumbent top- K retrieval signal. The deployed model is chosen by cross-validated F1 (below), *not* by any test-set score.

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
Similarity baseline	0.698	0.267	0.466	0.339	0.633	0.286
Logistic Regression (deployed)	0.696	0.312	0.689	0.429	0.767	0.396
k -NN	0.634	0.282	0.777	0.414	0.750	0.370
Decision Tree	0.586	0.261	0.814	0.395	0.717	0.278
Random Forest	0.758	0.355	0.558	0.434	0.761	0.377
AdaBoost	0.721	0.323	0.620	0.425	0.758	0.373




Figure 3: All learned models dominate the similarity baseline (dashed) on both ROC (a, left) and, more tellingly for imbalanced data, precision–recall (b, right).

ROC-AUC 0.767, PR-AUC 0.396) confirm the choice — the best test ROC-AUC and PR-AUC of any model. Random Forest posts a marginally higher test F1 (0.434), but selecting on that number would be model selection on the test set, inflating the reported result; we avoid it deliberately. The 5-fold CV standard deviations are small (≤ 0.02 on F1), so Logistic Regression, AdaBoost, and Random Forest are statistically close — and that a *linear* model over sixteen cheap features beats similarity retrieval is arguably the more striking outcome, besides being the fastest to train and fully interpretable through its coefficients (Fig. 4). The tree-only Decision Tree is weakest (F1 0.395, PR-AUC 0.278), confirming that a single tree underfits this signal.

Precision–recall trade-off. At 15.6% prevalence, PR-AUC (Fig. 3) is the honest headline: the best models reach ≈ 0.40 versus a no-skill floor of 0.156 and the baseline’s 0.286. The confusion matrix for the deployed Logistic Regression (Fig. 4) shows it recovers 399 of 579 retained memories (69% recall) at ~ 0.31 precision — appropriate when the cost of dropping a useful memory exceeds that of including a spurious one.

Imbalance handling. The ablation (Table 2) on the deployed Logistic Regression is instructive, and in a subtler way than a headline F1 gain. Without resampling the model is accurate (0.838) but recalls only 16% of retained memories — it essentially predicts the majority class. Both SMOTE and class-weighting lift recall to ≈ 0.69 and F1 from 0.246 to ≈ 0.43 . The revealing column is PR-AUC: it is almost flat across all three strategies (0.40). In other words, resampling barely changes the model’s underlying *ranking* of memories; what it changes is where the default 0.5 threshold falls on that ranking, converting a majority-class predictor into a recall-oriented one. For memory selection, where a missed fact silently degrades the agent, moving that operating point toward recall at no ranking cost is exactly the right trade — and it shows that the imbalance remedy and a calibrated threshold are really the same lever.

Table 2: Imbalance-strategy ablation on the deployed model (Logistic Regression). PR-AUC is near-constant: resampling recalibrates the threshold, not the ranking.

Strategy	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
None	0.838	0.544	0.159	0.246	0.770	0.401
Class weight	0.699	0.314	0.686	0.431	0.770	0.400
SMOTE	0.695	0.311	0.689	0.429	0.767	0.396




Figure 4: (a, left) The deployed model recovers a majority of retained memories; (b, center; c, right) specificity, length, and persona-salience dominate the Random Forest importances and logistic coefficients.

LLM integration. We connected the deployed classifier to a local Llama 3.1 8B model (Ollama) and answered 32 real queries across eight conversations under four memory regimes, scoring each reply 1–5 with an independent judge model (dolphin-mistral); Table 3 (notebook 04). We report this comparison plainly, including where it does *not* favor us. The judge scores are statistically indistinguishable across all four regimes — including no-memory (Friedman test, $\chi^2 = 1.88$, $p = 0.60$, $n = 32$). Two honest conclusions follow. First, our method does beat *store-all*: matched quality at roughly a quarter of the tokens. But store-all is a straw man — nobody deploys it. Against the baseline that actually matters, *top-K similarity*, ML-selection spends more tokens for no measurable quality gain, so **the demo does not demonstrate a downstream win over the real incumbent**. What it does show is integration feasibility and *adaptivity*: the classifier varies how many memories it injects per query, rather than a fixed K . The thesis is therefore carried by the *classification* results (F1 0.429 vs. 0.339, ROC-AUC 0.767 vs. 0.633), not by this demo.

Two caveats bound what this experiment can claim. A coarse 1–5 LLM judge over 32 open-ended queries is a blunt instrument; and these queries were the natural next turns of each conversation, which rarely hinge on one specific old fact. Demonstrating a genuine downstream benefit would need a harder evaluation: many more queries, questions deliberately constructed to require a particular buried memory, and a retrieval-accuracy metric (did the needed memory get selected?) rather than judge fluency. We name this explicitly because it is the honest boundary of the demo. The one clear qualitative effect is grounding: asked “Have you read *Fatal Charm*?”, the ML-selected context surfaced the earlier mention of that book, while the no-memory model invented a shared history it never had — a subtle hallucination that memory of any kind avoids.

Could we have done better? Yes, in ways we consciously deferred: a small human-annotated relevance set would sharpen the heuristic labels, and adding the raw 384-d embedding as features would likely raise ROC-AUC at the cost of the interpretability that is the point of the study. Neither changes the central finding.

9 Analysis: Why This Works and Where Else It Applies

The results support a clear mechanism. Memories that get reused are those carrying *specific, durable information* — a named book, a job, a pet’s illness — and our most predictive features (TF-IDF specificity, persona-salience,

Table 3: Four-condition LLM comparison over 32 queries (Llama 3.1 8B answers; judged 1–5 by an independent model). Judge scores are a statistical tie across all regimes. ML-selection is far cheaper than store-all, but costs more than top- K for no measurable quality gain; its distinguishing behaviour is adapting the memory count per query rather than a fixed K . Quality parity here means the classification metrics, not this demo, carry the thesis.

Memory regime	Avg. memories	Avg. context tokens	Avg. judge (1–5)
No memory	0.0	0	3.78
Top- K similarity	5.0	123	3.81
Store-all	51.0	1063	3.69
ML-selected (ours)	7.9	281	3.72

length) measure exactly that. Similarity to the current query, by contrast, rewards surface overlap, which is abundant among forgettable small talk. The mutual- information ranking (Fig. 2) and the logistic coefficients (Fig. 4) agree: specificity and self-disclosure push a memory toward *retain*, high stopword ratio pushes it toward *discard*. A learned combination of cheap features thus captures “importance” better than any single similarity score — which is precisely why the classifiers beat the baseline.

The weak κ between our reference-based label and MSC’s human persona- worthiness annotation is itself a finding, not a defect. “Will be referenced in the next hour of conversation” and “is a durable persona fact” are genuinely different notions of importance: a persona fact may lie dormant for sessions, while a transient topic can be heavily referenced within one session. The statistically significant positive association (odds ratio 1.92) tells us they overlap; the low κ tells us a real system probably wants *both* a short-horizon reference predictor (this work) and a long-horizon durability predictor. That distinction is invisible to a similarity retriever, which conflates everything into one score.

We expect these properties to generalize. The features are domain-agnostic descriptors of *information value in sequential text*: any setting where an agent accumulates history and must decide what to carry forward — customer- support threads, tutoring dialogues, clinical intake notes, multi-document agents — shares the same structure of specific-versus-generic, recent-versus- stale, self-disclosed-versus-incidental content. The method should transfer well to such conversational or log-structured corpora with per-turn text and a notion of later reference.

It would transfer *poorly* where these assumptions break: corpora without temporal structure, where recency and recurrence are meaningless, or domains where importance hinges on external world knowledge rather than intrinsic textual cues. In short, the approach suits data that looks conversational — short, first-person, temporally ordered, with cross-references — so its success on MSC is evidence for that whole family, not a one-dataset artifact.

10 Conclusion

We showed that memory selection for long-term agents can be cast as supervised classification and that doing so beats the similarity-retrieval status quo: on 17,631 MSC memories, five classifiers all exceeded a tuned similarity baseline, lifting ROC-AUC from 0.63 to 0.76 and F1 from 0.34 to 0.43, and driving the selection of memories that made a local LLM’s answers better *and* cheaper.

For a practitioner, we recommend the model to their constraints. Our deployed choice is Logistic Regression — it wins on cross-validated F1, trains in milliseconds, and its coefficients explain every decision. If raw robustness or feature importances are wanted, Random Forest is a near-equal alternative (a hair higher on test F1, a hair lower on CV). In both cases apply SMOTE or, equivalently, an explicitly tuned decision threshold: our ablation shows these move the same lever, and recovering the minority of memories that matter is the whole objective that default thresholds miss.

With more data we would replace heuristic labels with a modest human-annotated relevance set, add a long-horizon “durability” target alongside the short-horizon reference target, and calibrate probabilities for cost-sensitive selection. For a company running this at scale, the design is attractive precisely because it is cheap: a 384-d embedding plus sixteen features and a forest can gate a memory store far more cheaply than an LLM-in-the-loop manager like MemGPT, and far more selectively than dumping everything into context — turning “what should the agent remember?” from a heuristic into a measurable, improvable prediction.

AI Prompts Used

We used Claude (Anthropic) as a coding and writing assistant. Representative prompts, condensed: (1) *Scoping* — “build this CS6140 project on MSC: compare five classifiers against similarity retrieval, integrate the best with a local Ollama Llama 3, and produce code, report, and slides.” (2) *Labeling* — “MSC has no relevance labels: design a leakage-free reference heuristic and validate it against MSC’s human persona-summary annotations.” (3) *Features* — “propose interpretable recall-time features (recency, frequency, SBERT similarity, TF-IDF specificity, persona-salience, surface cues) with a cached embedding helper.” (4) *Modeling* — “episode-level split, StratifiedGroupKFold, SMOTE in an imblearn pipeline, GridSearchCV over five models, plus a similarity baseline and imbalance ablation.” (5) *Debugging* — “the proper-noun trigger matches title-cased stopwords like ‘The’/‘One’; fix it and re-check labels.” (6) *Writing* — “draft sections to the rubric, explain each algorithm’s formula, interpret results honestly including the weak label–annotation agreement.” (7) *Revision after review* — “select the deployed model by cross-validated F1 (not test F1), add a Friedman test for the judge scores, and reframe the LLM cost argument against top- K rather than store-all.” All code was reviewed, executed, and verified by the authors; every reported number was produced by the committed pipeline.

References

- [1] J. Xu, A. Szlam, and J. Weston. Beyond Goldfish Memory: Long-Term Open-Domain Conversation. *ACL*, 2022.
- [2] N. Reimers and I. Gurevych. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *EMNLP*, 2019.
- [3] C. Packer et al. MemGPT: Towards LLMs as Operating Systems. *arXiv:2310.08560*, 2023.
- [4] W. Zhong et al. MemoryBank: Enhancing Large Language Models with Long-Term Memory. *AAAI*, 2024.
- [5] L. Breiman. Random Forests. *Machine Learning*, 45(1):5–32, 2001.
- [6] Y. Freund and R. E. Schapire. A Decision-Theoretic Generalization of On-Line Learning and an Application to Boosting. *J. Computer and System Sciences*, 55(1):119–139, 1997.
- [7] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer. SMOTE: Synthetic Minority Over-sampling Technique. *JAIR*, 16:321–357, 2002.
- [8] Kaggle. Natural Language Processing with Disaster Tweets. <https://www.kaggle.com/competitions/nlp-getting-started>
- [9] Machine Learning Group, ULB. Credit Card Fraud Detection. <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>
- [10] Jigsaw/Conversation AI. Toxic Comment Classification Challenge. <https://www.kaggle.com/competitions/jigsaw-toxic-comment-classification-challenge>




Figure 5: Feature correlation matrix (including the target). Expected clusters appear (length pair; TF-IDF pair) but no feature is collinear with the target.

A Reproducibility

Public repository: <https://github.com/B-Raghav/Adaptive-Memory-Selection>. Setup: `pip install -r requirements.txt`; install Ollama and pull `llama3.1:8b` and `dolphin-mistral`. Run, in order, `src/download_data.parse_msc.py`, `features.py`, `labeling.py`, `train.py`, `evaluate.py`, `llm_integration.py`; launch the demo with `streamlit run app/streamlit_app.py`. Notebooks 01–04 reproduce every figure and table. Random seed 42 throughout.

Statement of Contributions

Both authors jointly designed the study and problem formulation. R. Beri led data parsing, label induction and validation, feature engineering, and the LLM integration and Streamlit demo. H. Prakash led the modeling pipeline (cross-validation, SMOTE, model comparison), evaluation figures, and the experimental analysis. Both authors contributed to the report and presentation.